

churnproj

July 28, 2023

Author: Ashley Eastman

Spring 2023 - Customer Churn Prediction Analysis

```
[1]: # Libraries
# General
import pandas as pd
import numpy as np

# ML
from sklearn.model_selection import train_test_split
import statsmodels.formula.api as smf
from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score, roc_curve, auc
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.preprocessing import OneHotEncoder

# Viz
import matplotlib.pyplot as plt
from importlib import *
import plotly.express as px
import matplotlib.patches as patches
```

This dataset includes information about cellphone company customers and whether they stayed or left the company (churn)

Note: Data Cleaning has been performed prior to this step

```
[2]: # Dataset: customerchurn.csv
churn = pd.read_csv("customerchurn.csv")
churn.info()
churn
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Churn                  7032 non-null   int64
```

```

1 MonthlyCharges 7032 non-null float64
2 SeniorCitizen 7032 non-null int64
3 PaymentMethod 7032 non-null object
4 InternetService 7032 non-null object
5 tenure 7032 non-null int64
6 Contract 7032 non-null object
dtypes: float64(1), int64(3), object(3)
memory usage: 384.7+ KB

```

```

[2]:      Churn  MonthlyCharges  SeniorCitizen  PaymentMethod  InternetService \
0         0         29.85             0  Electronic check          DSL
1         0         56.95             0    Mailed check          DSL
2         1         53.85             0    Mailed check          DSL
3         0         42.30             0    Bank transfer          DSL
4         1         70.70             0  Electronic check  Fiber optic
...     ...         ...             ...         ...         ...
7027      0         84.80             0    Mailed check          DSL
7028      0        103.20             0    Credit card  Fiber optic
7029      0         29.60             0  Electronic check          DSL
7030      1         74.40             1    Mailed check  Fiber optic
7031      0        105.65             0    Bank transfer  Fiber optic

      tenure      Contract
0         1  Month-to-month
1        34    One year
2         2  Month-to-month
3        45    One year
4         2  Month-to-month
...     ...         ...
7027      24    One year
7028      72    One year
7029      11  Month-to-month
7030       4  Month-to-month
7031      66    Two year

```

[7032 rows x 7 columns]

```

[3]: # Dependent variable: Churn 0 (no churn) or 1 (churn)
# Independent variables: 6
X = churn.drop(['Churn'], axis=1)
y = churn['Churn']
print(X.shape), print(y.shape)

```

```

(7032, 6)
(7032,)

```

```

[3]: (None, None)

```

0.0.1 EDA

```
[4]: # some EDA

# Null value check
nullvals = churn.isnull().sum()
print('Null values in each col:\n', nullvals)

f = plt.figure(figsize=(8, 6))
# Creating a gradient colormap based on 'Pastel1'
pastel1 = plt.get_cmap('Pastel1')
n = len(churn.columns)
gradient_colors = pastel1(np.linspace(0, 1, n**2))

# Create a new colormap from the gradient colors
gradient_cmap = plt.matplotlib.colors.LinearSegmentedColormap.
    ↪from_list('pastel1_gradient', gradient_colors, n**2)

plt.matshow(churn.corr(), fignum=f.number, cmap=gradient_cmap)
plt.xticks(range(churn.select_dtypes(['number']).shape[1]), churn.
    ↪select_dtypes(['number']).columns, fontsize=14, rotation=45)
plt.yticks(range(churn.select_dtypes(['number']).shape[1]), churn.
    ↪select_dtypes(['number']).columns, fontsize=14)
cb = plt.colorbar()
cb.ax.tick_params(labelsize=8)

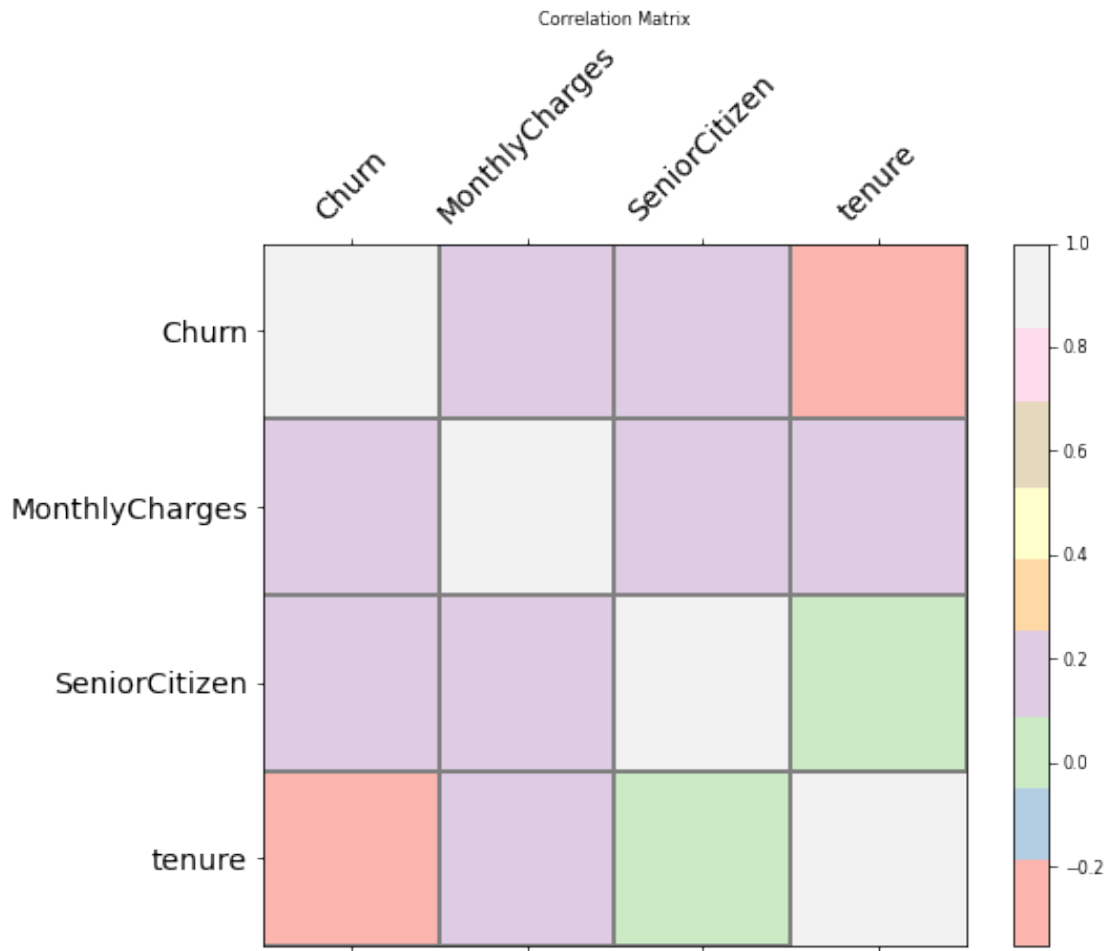
# Black Borders
for i in range(len(churn.corr())):
    for j in range(len(churn.corr())):
        if i != j:
            plt.gca().add_patch(patches.Rectangle((i - 0.5, j - 0.5), 1, 1,
    ↪fill=False, edgecolor='gray', lw=2))

plt.title('Correlation Matrix', fontsize=8)
plt.show()
```

Null values in each col:

Churn	0
MonthlyCharges	0
SeniorCitizen	0
PaymentMethod	0
InternetService	0
tenure	0
Contract	0


dtype: int64



```
[5]: # Using seaborn to plot scatter plot for continuous variables
import seaborn as sns
sns.pairplot(churn, diag_kind='kde', height=2.6, hue='Churn', palette='bwr',
             plot_kws=dict(s=10))
```

```
/opt/homebrew/lib/python3.9/site-packages/seaborn/axisgrid.py:118: UserWarning:
The figure layout has changed to tight
    self._figure.tight_layout(*args, **kwargs)
```

```
[5]: <seaborn.axisgrid.PairGrid at 0x1668afd60>
```



```
[6]: color_map = {1: '#FFB3B3', 0: '#D2DAFF'}
fig = px.histogram(churn, x="Churn", color="SeniorCitizen",
    title="<b>Distribution Chart of Senior Citizens & Churn Rate</b>",
    color_discrete_map=color_map)
fig.update_layout(width=500, height=400, bargap=0.1)
fig.update_traces(marker_line_color='gray', marker_line_width=1.5)
fig.show()
```

```
[7]: senior_churn = churn[churn['SeniorCitizen'] == 1]
senior = senior_churn['Churn'].value_counts()
senior[1] / (senior[0] + senior[1])
seniorYes = senior_churn[senior_churn['Churn'] == 1]
seniorNo = senior_churn[senior_churn['Churn'] == 0]
total = len(seniorYes) / (len(seniorYes) + len(seniorNo))
```

```
print('In total, approximately: ', round(total, 2) * 100, ' % of senior_
citizens churned')
```

In total, approximately: 42.0 % of senior citizens churned

```
[8]: show = churn[churn['Churn'] == 1]
seniorchurn = show[show['SeniorCitizen'] == 1]
seniorchurn
```

```
[8]:      Churn  MonthlyCharges  SeniorCitizen  PaymentMethod  InternetService \
20         1          39.65             1  Electronic check          DSL
53         1          80.65             1      Credit card  Fiber optic
55         1          95.45             1  Electronic check  Fiber optic
99         1          98.50             1  Electronic check  Fiber optic
113        1          76.50             1  Electronic check  Fiber optic
...      ...              ...              ...              ...
6982        1          88.05             1  Electronic check  Fiber optic
6997        1          75.05             1      Credit card  Fiber optic
6999        1          74.45             1  Electronic check  Fiber optic
7021        1          75.75             1  Electronic check  Fiber optic
7030        1          74.40             1      Mailed check  Fiber optic

      tenure      Contract
20         1  Month-to-month
53         8  Month-to-month
55        18  Month-to-month
99        25  Month-to-month
113       37  Month-to-month
...      ...              ...
6982       50  Month-to-month
6997        3  Month-to-month
6999        1  Month-to-month
7021        1  Month-to-month
7030        4  Month-to-month
```

[476 rows x 7 columns]

Notes on above EDA:

- We can see that monthly charges are not a huge identifier of a customer churning or not - as we have many customers with high monthly charges that have not churned.
- Of the 1,393 customers that churned, 476 of them were senior citizens.
- Of the 1,142 seniors in the dataset, 476 of them churned which is a very high churn rate for senior citizens. (41%)
- Tenure has a negative correlation with churn as we can also see from both the correlation matrix above & our pairplot
- When Tenure is low, and monthly charges are high, we can visually identify a high churn rate

```
[9]: # Splitting train/test
train, test = train_test_split(churn, test_size=0.3, random_state=42)
train.shape, test.shape
```

```
[9]: ((4922, 7), (2110, 7))
```

Baseline Accuracy - Predict the most common output feature

```
[10]: baseline = y.value_counts()[0]/len(y)
print('Baseline Accuracy:', round(baseline, 3) * 100, '%')
```

Baseline Accuracy: 73.4 %

```
[11]: # Variable Selecting & model fitting
logreg = smf.logit(formula = 'Churn ~ MonthlyCharges + SeniorCitizen + 
    ↳PaymentMethod + InternetService + tenure + Contract',
                    data = train).fit()
print(logreg.summary())
```

Optimization terminated successfully.

Current function value: 0.421207

Iterations 8

Logit Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                  Logit    Df Residuals:              4911
Method:                  MLE     Df Model:                  10
Date:                   Fri, 28 Jul 2023    Pseudo R-squ.:          0.2725
Time:                   17:56:54    Log-Likelihood:         -2073.2
converged:               True     LL-Null:                 -2849.7
Covariance Type:         nonrobust    LLR p-value:            0.000
=====
=====
```

	coef	std err	z	P> z
[0.025 0.975]				
Intercept	-0.6747	0.214	-3.157	0.002
-1.094 -0.256				
PaymentMethod[T.Credit card]	-0.1505	0.135	-1.115	0.265
-0.415 0.114				
PaymentMethod[T.Electronic check]	0.3186	0.111	2.865	0.004
0.101 0.537				
PaymentMethod[T.Mailed check]	-0.1227	0.134	-0.915	0.360
-0.385 0.140				
InternetService[T.Fiber optic]	0.9479	0.156	6.059	0.000
0.641 1.255				
InternetService[T.No]	-0.7120	0.182	-3.910	0.000

-1.069	-0.355				
Contract[T.One year]		-0.8964	0.128	-7.006	0.000
-1.147	-0.646				
Contract[T.Two year]		-1.5712	0.202	-7.760	0.000
-1.968	-1.174				
MonthlyCharges		0.0048	0.004	1.293	0.196
-0.002	0.012				
SeniorCitizen		0.3610	0.098	3.677	0.000
0.169	0.553				
tenure		-0.0322	0.003	-12.417	0.000
-0.037	-0.027				

=====

=====

Notes / Interpretation on Logit Regression Results:

- Our $p > |z|$ * value for payment method is large (0.265, 0.360)
- So I am going to create a column that only includes Electronic Check as payment method as Credit card & mailed check can be contributing to high variance inflation factors and should therefore be removed for a better fit model.

```
[12]: # Remove PaymentMethod Credit Card & Payment Method Mailed check.
# Create a new feature (dummy var)
train2 = train.copy()
train2['ElectronicCheck'] = (train2['PaymentMethod'] == 'Electronic check').
    ↪astype('int64')
train2.drop(columns=['PaymentMethod'], inplace=True)

# Repeat for test set as well
test2 = test.copy()
test2['ElectronicCheck'] = (test2['PaymentMethod'] == 'Electronic check').
    ↪astype('int64')
test2.drop(columns=['PaymentMethod'], inplace=True)

test2.head()
```

```
[12]:      Churn  MonthlyCharges  SeniorCitizen  InternetService  tenure  \
2476      0           25.00             1             No         61
6773      0           24.70             0             No         19
6116      1          102.25             0  Fiber optic         13
3047      0           55.05             0             DSL         37
4092      0           29.45             0             DSL          6

      Contract  ElectronicCheck
2476      Two year              0
6773  Month-to-month              0
6116  Month-to-month              0
3047  Month-to-month              0
```


4092 Month-to-month 0

```
[13]: test2['ElectronicCheck'].value_counts()
```

```
[13]: 0    1418
      1     692
      Name: ElectronicCheck, dtype: int64
```

```
[14]: logreg2 = smf.logit(formula = 'Churn ~ SeniorCitizen + ElectronicCheck +
      ↪InternetService + tenure + Contract',
      data = train2).fit()
      print(logreg2.summary())
```

Optimization terminated successfully.

Current function value: 0.421529

Iterations 8

Logit Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          4922
Model:                Logit    Df Residuals:              4914
Method:                MLE     Df Model:                  7
Date:                  Fri, 28 Jul 2023    Pseudo R-squ.:        0.2719
Time:                  17:56:55    Log-Likelihood:         -2074.8
converged:              True     LL-Null:              -2849.7
Covariance Type:        nonrobust    LLR p-value:            0.000
=====
```

```
=====
                                coef    std err          z      P>|z|
-----
[0.025    0.975]
-----
Intercept                   -0.5498     0.085    -6.465     0.000
-0.716    -0.383
InternetService[T.Fiber optic]  1.1192     0.090    12.471     0.000
0.943     1.295
InternetService[T.No]        -0.8659     0.141    -6.120     0.000
-1.143    -0.589
Contract[T.One year]        -0.8797     0.127    -6.936     0.000
-1.128    -0.631
Contract[T.Two year]        -1.5495     0.201    -7.703     0.000
-1.944    -1.155
SeniorCitizen                0.3587     0.098     3.656     0.000
0.166     0.551
ElectronicCheck              0.4130     0.081     5.100     0.000
0.254     0.572
tenure                      -0.0308     0.002   -12.923     0.000
-0.035    -0.026
```

=====

Notes on above Logit Regression results: - Now we can see that all of the p-values are close to 0 now and we can safely assess that we have dropped all unimportant variables from the dataset

Predicting Probability of Customer Churn & Deciding the threshold value

Below I have compared multiple thresholds to determine the optimal p^* to reduce error rate in our predictor Side note: Using optimal bayes formula

0.0.2 Conditional Probability (optimal bayes)

Formula: $f(x) = P(Y = 1 \mid X = x)$

```
[15]: # 1. Predicting the probability of default
# y_prob will be a vector of probabilities that the customer will churn or not
# (1 or 0)
y_prob = logreg2.predict(test2)

# 2. Determining the optimal threshold of the default probability
## The threshold of high churn-risk and low churn-risk will be 0.333
## If  $p > 1/3$ , then we consider high churn-risk and should send a promotion to
# the customer or check in with them

# 3. Predicting the label
y_pred = pd.Series([1 if x > 1/3 else 0 for x in y_prob], index=y_prob.index)

# use  $P = 1/2$  as the threshold
y_pred_2 = pd.Series([1 if x > 1/2 else 0 for x in y_prob], index=y_prob.index)
```

P value = $1/3$

```
[16]: # Now we have our probability of default, I will construct a confusion matrix
# based on decision tree threshold we have computed
# High churn risk > 0.333.
# or high churn risk

y_test = test2['Churn']

cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix : \n", cm)

acc = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)
print('Classification Report: ', report)
print('Accuracy: ', acc)
```

Confusion Matrix :

```
[[1199 350]
```

```
[ 145 416]]
```

Classification Report:

			precision	recall	f1-score	support
	0	0.89	0.77	0.83		1549
	1	0.54	0.74	0.63		561
	accuracy		0.77			2110
	macro avg	0.72	0.76	0.73		2110
	weighted avg	0.80	0.77	0.78		2110

Accuracy: 0.7654028436018957

With our threshold at $P = 1/2$, we can see that recall has increased from ~77% -> ~89% for predicting churn while f1-score has increased from 83% -> ~86%. So depending on if we care more about correctly identifying true positives or true negatives, vs incorrectly identifying them (false positives/false negatives) then we should use 1/2 as our optimal p-value threshold.

```
[17]: # Finding optimal p value threshold for churn prediction
y_test2 = test2['Churn']

cm2 = confusion_matrix(y_test2, y_pred_2)
print("Confusion Matrix : \n", cm2)

acc2 = accuracy_score(y_test2, y_pred_2)
report2 = classification_report(y_test2, y_pred_2)
print('Classification Report: ', report2)
print('Accuracy: ', acc2)
```

Confusion Matrix :

```
[[1382 167]
```

```
[ 279 282]]
```

Classification Report:

			precision	recall	f1-score	support
	0	0.83	0.89	0.86		1549
	1	0.63	0.50	0.56		561
	accuracy		0.79			2110
	macro avg	0.73	0.70	0.71		2110
	weighted avg	0.78	0.79	0.78		2110

Accuracy: 0.7886255924170616

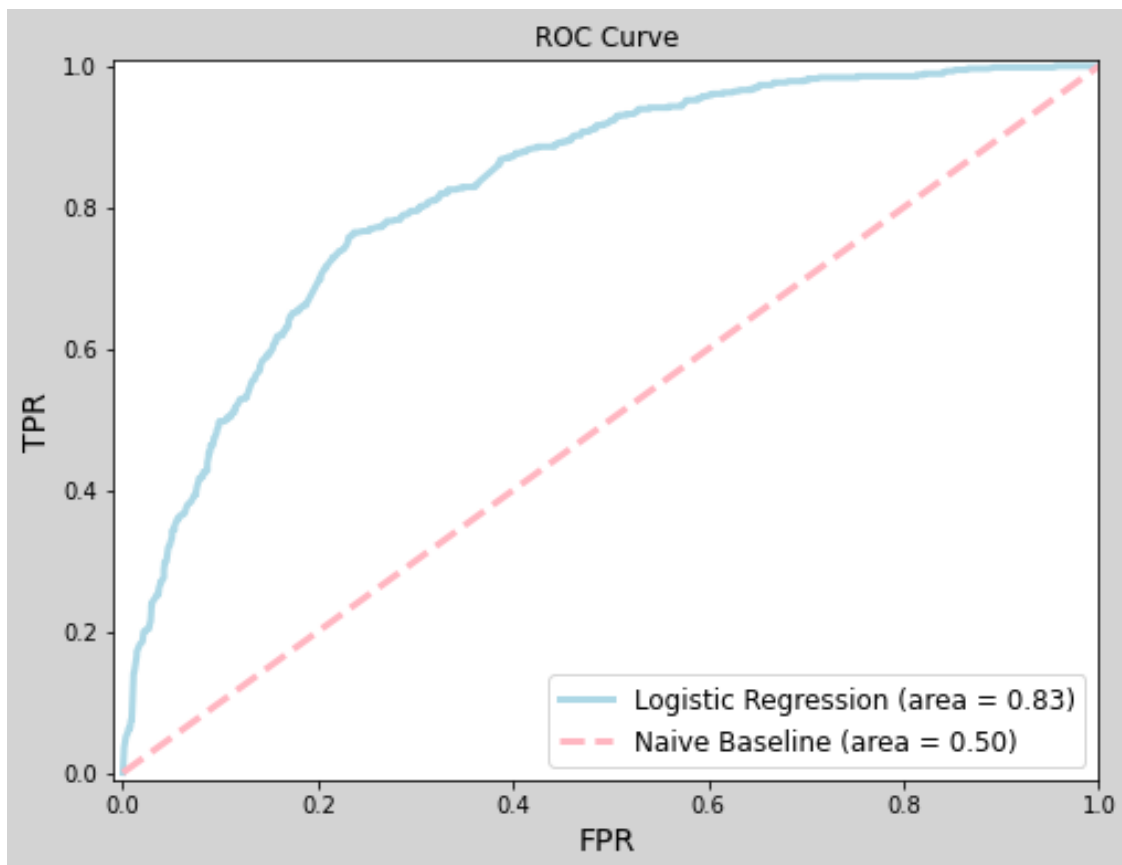
Precision: 83% The ROC curve plots the TPR and FPR for every break-even threshold p between 0.0 and 1.0

```
[18]: y_train = train2['Churn']
x_train = train2.drop(['Churn'], axis=1)

y_test = test2['Churn']
x_test = test2.drop(['Churn'], axis=1)

fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6), facecolor='lightgray')
plt.title('ROC Curve', fontsize=12)
plt.xlabel('FPR', fontsize=14)
plt.ylabel('TPR', fontsize=14)
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr, tpr, lw=3, color = 'lightblue', label='Logistic Regression (area_
    ↳{:0.2f})'.format(roc_auc))
plt.plot([0, 1], [0, 1], color='lightpink', lw=3, linestyle='--', label='Naive_
    ↳Baseline (area = 0.50)')
plt.legend(loc='lower right', fontsize=12)
plt.show()
```



LDA Model

```
[19]: # First I will change categorical variables to binary variables using OHE
```

```
# initialize the OneHotEncoder
drop_enc = OneHotEncoder(drop='first').
    ↳fit(x_train[['InternetService', 'Contract']])
print(drop_enc.categories_)
```

```
[array(['DSL', 'Fiber optic', 'No'], dtype=object), array(['Month-to-month',
'One year', 'Two year'], dtype=object)]
```

```
[20]: # Perform the transformation for both the training and the test set.
```

```
X_train_categorical = drop_enc.
    ↳transform(x_train[['InternetService', 'Contract']]).toarray()
X_train_numerical = x_train[['MonthlyCharges', 'SeniorCitizen', 'tenure']].values
# combine the numerical variables and the one-hot encoded categorical variables
X_train_transformed = np.concatenate((X_train_numerical, X_train_categorical),
    ↳axis = 1)

X_test_categorical = drop_enc.transform(x_test[['InternetService', 'Contract']]).
    ↳toarray()
X_test_numerical = x_test[['MonthlyCharges', 'SeniorCitizen', 'tenure']].values
X_test_transformed = np.concatenate((X_test_numerical, X_test_categorical), axis=
    ↳1)
```

```
[21]: lda = LinearDiscriminantAnalysis()
lda.fit(X_train_transformed, y_train)

y_prob_lda = lda.predict_proba(X_test_transformed)
y_pred_lda = pd.Series([1 if x > 1/2 else 0 for x in y_prob_lda[:,1]],
    ↳index=y_prob.index)

cm = confusion_matrix(y_test, y_pred_lda)
print ("Confusion Matrix: \n", cm)
print ("\nAccuracy:", accuracy_score(y_test, y_pred_lda))
```

Confusion Matrix:

```
[[1365  184]
 [ 268  293]]
```


Accuracy: 0.785781990521327

```
[22]: fpr_lda, tpr_lda, _ = roc_curve(y_test, y_prob_lda[:,1])
roc_auc_lda = auc(fpr_lda, tpr_lda)
```

```

plt.figure(figsize=(8, 6), facecolor='lightgray')
plt.title('ROC Curve', fontsize=14)
plt.xlabel('FPR', fontsize=12)
plt.ylabel('TPR', fontsize=12)
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr, tpr, lw=3, color = 'pink', label='Logistic Regression (area = {:.
    ↪2f})'.format(roc_auc))
plt.plot(fpr_lda, tpr_lda, color = 'lightblue', lw=3, label='LDA (area = {:.
    ↪2f})'.format(roc_auc_lda))
plt.plot([0, 1], [0, 1], color='darkgray', lw=3, linestyle='--', label='Naive_
    ↪Baseline (area = 0.50)')
plt.legend(loc='lower right', fontsize=14)
plt.show()

```



Feature Importance for LDA

```
[23]: from sklearn.ensemble import RandomForestClassifier
```

```

X_train_df = pd.DataFrame(X_train_transformed, columns=train2.columns)
X_test_df = pd.DataFrame(X_test_transformed, columns=test2.columns)

rf = RandomForestClassifier(n_estimators=200, random_state=42)
rf.fit(X_train_transformed, y_train)
y_pred = rf.predict(X_test_transformed)
score = rf.score(X_test_transformed, y_test)
important_feats = pd.Series(rf.feature_importances_, index=X_train_df.columns).
    ↪sort_values(ascending=False)
print(important_feats)

```

```

Churn          0.513605
SeniorCitizen  0.307479
InternetService 0.056721
ElectronicCheck 0.052663
Contract       0.034173
tenure         0.020344
MonthlyCharges 0.015014
dtype: float64

```

```

[24]: # Visualizing Feature Importance
importances = rf.feature_importances_
weights = pd.Series(importances, index=X_train_df.columns)
weights = weights.drop(['Churn'])


# Sort the 'weights' Series in ascending order
weights_sorted = weights.sort_values(ascending=False)

# Generate a list of different colors for each bar
num_features = len(weights)
colors = plt.cm.Pastel2(np.linspace(0, 1, num_features))

# Now, create the sorted bar plot using the 'plot()' method and set the 'color' ↪
    ↪parameter
weights_sorted.plot(kind='bar', color=colors, edgecolor = 'gray', figsize=(8, ↪
    ↪6))

# Optionally, you can add labels and title
plt.xlabel('Features')
plt.ylabel('Importance')
plt.title('Feature Importance (Sorted)')
plt.show()

```



Conclusion & Notes

- With the data that I have available, I have observed that our most important feature to determine churn rate is if the customer is a senior citizen or not (most senior citizens in this dataset churn)
- Lowered variance inflation factors by removing PaymentMethod Credit Card & Payment Method Mailed check for a more a better fit model with p values $\rightarrow 0$
- I minimized the error rate by finding the optimal p value threshold (1/2)
- With more data on the customers, a better accuracy rate and determination of churn rate factors can be determined. For example, if I had access to which platform the customers were using (app, mobile web) or more specifics on the types of services the customer used most, or even if they had dependents then we would be able and built a better model for determining churn

Additional Notes on EDA:

- We can see that monthly charges are not a huge identifier of a customer churning or not - as we have many customers with high monthly charges that have not churned.
- Of the 1,393 customers that churned, 476 of them were senior citizens.
- Of the 1,142 seniors in the dataset, 476 of them churned which is a very high churn rate for senior citizens. (41%)
- Tenure has a negative correlation with churn as we can also see from both the correlation matrix above & our pairplot
- When Tenure is low, and monthly charges are high, we can visually identify a high churn rate

ashleyha@berkeley.edu